

# Uncovering Sentiments using EDGAR Datasets

Team 9

Dawna Grace Raj

Jai Soni

Nikhil Kohli

## Experiments:

### Experiment 1:

#### Loading the data


| sentiment | text  |
|-----------|---|
| positive  | thank good afternoon everyone welcome nvidias ... |
| positive  | look past q1 expect channel inventory correcti... |
| negative  | china game weakness slow economic environment ... |
| negative  | dont know could tear apart tease apart harlan ... |
| positive  | thank ill turn call back jenhsun close remark     |

## Bag of Words Using Count Vectorizer

```
[ ] from sklearn.feature_extraction.text import CountVectorizer
```

```
[ ] vectorizer = CountVectorizer(min_df=0, lowercase=False)
```

```
[ ] vectorizer.fit(new_df['text'])
```

 CountVectorizer(analyzer='word', binary=False, decode\_error='strict', dtype=<class 'numpy.int64'>, encoding='utf-8', input='content', lowercase=False, max\_df=1.0, max\_features=None, min\_df=0, ngram\_range=(1, 1), preprocessor=None, stop\_words=None, strip\_accents=None, token\_pattern='(?u)\\b\\w\\w+\\b', tokenizer=None, vocabulary=None)

### ➤ Converting it to an array of integers

```
[ ] a=vectorizer.transform(new_df['text']).toarray()
```

```
[ ] a
```

```
array([[0, 0, 0, ..., 0, 0, 0],  
       [0, 0, 0, ..., 0, 0, 0],  
       [0, 0, 0, ..., 0, 0, 0],  
       ...,  
       [0, 0, 0, ..., 0, 0, 0],  
       [0, 0, 0, ..., 0, 0, 0],  
       [0, 0, 0, ..., 0, 0, 0]], dtype=int64)
```

### ➤ Creating Train and Test Data

```
[ ] from sklearn.model_selection import train_test_split
```

```
[ ] X_train, X_test, y_train, y_test = train_test_split(a, Y, test_size=0.2, random_state=1000)
```

#### Creating Train and Test Data

```
[ ] from sklearn.model_selection import train_test_split
```

```
[ ] X_train, X_test, y_train, y_test = train_test_split(a, Y, test_size=0.2, random_state=1000)
```

## Creating a Logistic Regression Model:

```
[ ] from sklearn.linear_model import LogisticRegression
```

```
[ ] classifier = LogisticRegression()  
classifier.fit(X_train, y_train)  
score = classifier.score(X_test, y_test)
```

```
[ ] print("Accuracy:", score)
```



Accuracy: 0.8098159509202454

Obtained an Accuracy of 81%

```
[ ] y_pred = classifier.predict(X_test)  
y_pred
```



```
array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 0, 0,  
       1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
       1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1,  
       1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1,  
       1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
       1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
       1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1,  
       1, 1, 1, 1, 1, 1, 1, 1, 1, 1], dtype=int64)
```

## Creating a metrics to measure model performance

```
from sklearn.metrics import confusion_matrix
confusion_matrix(y_test, y_pred)
```

```
array([[ 5, 24],
       [ 7, 127]], dtype=int64)
```


```
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

|             | precision | recall | f1-score | support |
|-------------|-----------|--------|----------|---------|
| 0           | 0.42      | 0.17   | 0.24     | 29      |
| 1           | 0.84      | 0.95   | 0.89     | 134     |
| avg / total | 0.77      | 0.81   | 0.78     | 163     |

```
[ ] import matplotlib.pyplot as plt
import seaborn as sns

ax= plt.subplot()
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, ax = ax, fmt='g'); #annot=True to annotate cells

# labels, title and ticks
ax.set_xlabel('Predicted labels');ax.set_ylabel('True labels');
ax.set_title('Confusion Matrix');
ax.xaxis.set_ticklabels(['Negative', 'Positive']); ax.yaxis.set_ticklabels(['Positive', 'Negative']);
```



## Fully connected Model using Keras

```
model_bow.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
model_bow.summary()
```

| Layer (type)        | Output Shape | Param # |
|---------------------|--------------|---------|
| dense_10 (Dense)    | (None, 32)   | 128032  |
| dropout_7 (Dropout) | (None, 32)   | 0       |
| dense_11 (Dense)    | (None, 32)   | 1056    |
| dropout_8 (Dropout) | (None, 32)   | 0       |
| dense_12 (Dense)    | (None, 1)    | 33      |

=====  
Total params: 129,121  
Trainable params: 129,121  
Non-trainable params: 0

```
history_bow = model_bow.fit(X_train, y_train, epochs=5, verbose=False, validation_data=(X_test, y_test), batch_size=10)
```



```
loss, accuracy = model_bow.evaluate(X_train, y_train)
print("Training Accuracy: {:.4f}".format(accuracy))
loss, accuracy = model_bow.evaluate(X_test, y_test)
print("Testing Accuracy: {:.4f}".format(accuracy))
```


```
648/648 [=====] - 0s 54us/step
Training Accuracy: 0.8673
163/163 [=====] - 0s 61us/step
Testing Accuracy: 0.8282
```

Test Loss: 0.4681244474247189

Test Accuracy 0.8282208592613782

```
dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
```





## Creating a LSTM model using keras

```
[ ] model_lstm = Sequential()
    model_lstm.add(Embedding(20000, 100, input_length=50))
    model_lstm.add(LSTM(100,activation='relu', dropout=0.5))
    model_lstm.add(Dense(1, activation='sigmoid'))
    model_lstm.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

Train on 648 samples, validate on 163 samples

Epoch 1/4

648/648 [=====] - 5s 8ms/step - loss: 0.6542 - acc: 0.7608 - val\_loss: 0.5946 - val\_acc: 0.7485

Epoch 2/4

648/648 [=====] - 2s 3ms/step - loss: 0.5361 - acc: 0.8210 - val\_loss: 0.5504 - val\_acc: 0.7485

Epoch 3/4

648/648 [=====] - 2s 3ms/step - loss: 0.4030 - acc: 0.8210 - val\_loss: 0.5443 - val\_acc: 0.7485

Epoch 4/4



648/648 [=====] - 2s 3ms/step - loss: 0.3345 - acc: 0.8210 - val\_loss: 0.5319 - val\_acc: 0.7485



Test Loss: 0.4681244474247189

Test Accuracy 0.8282208592613782

```
dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
```



**Creating a Glove model**

```

sequence_input = Input(shape=(50,), dtype='int32')
embedded_sequences = embedding_layer(sequence_input)
x = Conv1D(128, 5, activation='relu')(embedded_sequences)
x = MaxPooling1D(5)(x)
x = Conv1D(128, 5, activation='relu')(x)
x = MaxPooling1D(5)(x)
x = Conv1D(128, 5, activation='relu', data_format='channels_first')(x)
x = MaxPooling1D(35)(x) # global max pooling
x = Flatten()(x)
x = Dense(128, activation='relu')(x)
preds = Dense(1, activation='softmax')(x)

model_glove = Model(sequence_input, preds)
model_glove.compile(loss='binary_crossentropy',
                    optimizer='adam',
                    metrics=['acc'])

# happy learning!
history_glove=model_glove.fit(x_train, y_train, validation_data=(x_val, y_val), epochs=10, batch_size=32)

```

Train on 649 samples, validate on 162 samples

```

Epoch 1/10
649/649 [=====] - 4s 6ms/step - loss: 2.9477 - acc: 0.8151 - val_loss: 3.6412 - val_acc: 0.7716
Epoch 2/10
649/649 [=====] - 1s 1ms/step - loss: 2.9477 - acc: 0.8151 - val_loss: 3.6412 - val_acc: 0.7716
Epoch 3/10
649/649 [=====] - 1s 1ms/step - loss: 2.9477 - acc: 0.8151 - val_loss: 3.6412 - val_acc: 0.7716
Epoch 4/10
649/649 [=====] - 1s 1ms/step - loss: 2.9477 - acc: 0.8151 - val_loss: 3.6412 - val_acc: 0.7716
Epoch 5/10
649/649 [=====] - 1s 1ms/step - loss: 2.9477 - acc: 0.8151 - val_loss: 3.6412 - val_acc: 0.7716
Epoch 6/10
649/649 [=====] - 1s 1ms/step - loss: 2.9477 - acc: 0.8151 - val_loss: 3.6412 - val_acc: 0.7716
Epoch 7/10
649/649 [=====] - 1s 1ms/step - loss: 2.9477 - acc: 0.8151 - val_loss: 3.6412 - val_acc: 0.7716
Epoch 8/10
649/649 [=====] - 1s 1ms/step - loss: 2.9477 - acc: 0.8151 - val_loss: 3.6412 - val_acc: 0.7716
Epoch 9/10
649/649 [=====] - 1s 1ms/step - loss: 2.9477 - acc: 0.8151 - val_loss: 3.6412 - val_acc: 0.7716
Epoch 10/10
649/649 [=====] - 1s 1ms/step - loss: 2.9477 - acc: 0.8151 - val_loss: 3.6412 - val_acc: 0.7716

```

## Getting a Test accuracy of 0.7716 with the Glove model

### Trying out different models to get better accuracy

#### Model 2

```

from keras.optimizers import adam
sequence_input = Input(shape=(50,), dtype='int32')
embedded_sequences = embedding_layer_new(sequence_input)
x = Conv1D(32, 5, activation='relu')(embedded_sequences)
x = MaxPooling1D(5)(x)
#x = Conv1D(64, 5, activation='relu', data_format='channels_first')(x)
#x = MaxPooling1D(5)(x) # global max pooling
x = Flatten()(x)
x = Dense(32, activation='relu')(x)
preds = Dense(1, activation='sigmoid')(x)



model_glove2 = Model(sequence_input, preds)
model_glove2.compile(loss='binary_crossentropy',
                    optimizer=adam(lr=0.00001),
                    metrics=['accuracy'])

# happy learning!
history_glove2=model_glove2.fit(x_train, y_train, validation_data=(x_val, y_val), epochs=10, batch_size=32)

```

Train on 649 samples, validate on 162 samples

```
Epoch 1/10
649/649 [=====] - 3s 5ms/step - loss: 0.5295 - acc: 0.8136 - val_loss: 0.5935 - val_acc: 0.7778
Epoch 2/10
649/649 [=====] - 0s 413us/step - loss: 0.5260 - acc: 0.8136 - val_loss: 0.5919 - val_acc: 0.7778
Epoch 3/10
649/649 [=====] - 0s 421us/step - loss: 0.5235 - acc: 0.8136 - val_loss: 0.5905 - val_acc: 0.7778
Epoch 4/10
649/649 [=====] - 0s 418us/step - loss: 0.5209 - acc: 0.8151 - val_loss: 0.5892 - val_acc: 0.7778
Epoch 5/10
649/649 [=====] - 0s 438us/step - loss: 0.5185 - acc: 0.8151 - val_loss: 0.5884 - val_acc: 0.7778
Epoch 6/10
649/649 [=====] - 0s 421us/step - loss: 0.5161 - acc: 0.8151 - val_loss: 0.5876 - val_acc: 0.7778
Epoch 7/10
649/649 [=====] - 0s 433us/step - loss: 0.5139 - acc: 0.8151 - val_loss: 0.5866 - val_acc: 0.7778
Epoch 8/10
649/649 [=====] - 0s 422us/step - loss: 0.5119 - acc: 0.8151 - val_loss: 0.5856 - val_acc: 0.7778
Epoch 9/10
649/649 [=====] - 0s 420us/step - loss: 0.5099 - acc: 0.8166 - val_loss: 0.5841 - val_acc: 0.7778
Epoch 10/10
649/649 [=====] - 0s 427us/step - loss: 0.5081 - acc: 0.8166 - val_loss: 0.5829 - val_acc: 0.7778
```



Model 3:

## Model 3 Trainable True Plot


```
[ ] import matplotlib.pyplot as plt
    #print(metrics.accuracy_score(Y_test, Y_predicted))

    #score = model_lstm.evaluate(X_test, y_test, verbose=3)
    #print('Test Loss: ', score[0])
    #print('Test Accuracy', score[1])

    # list all data in history
    print(history_glove2.history.keys())
    # summarize history for accuracy
    plt.plot(history_glove2.history['acc'])
    plt.plot(history_glove2.history['val_acc'])
    plt.title('Model Accuracy')
    plt.ylabel('accuracy')
    plt.xlabel('epoch')
    plt.legend(['train', 'test'], loc='upper left')
    plt.show()
    # summarize history for loss
    plt.plot(history_glove2.history['loss'])
    plt.plot(history_glove2.history['val_loss'])
    plt.title('Model Loss')
    plt.ylabel('loss')
    plt.xlabel('epoch')
    plt.legend(['train', 'test'], loc='upper left')
    plt.show()
```

Train on 649 samples, validate on 162 samples

```
Epoch 1/10
649/649 [=====] - 4s 6ms/step - loss: 0.5631 - acc: 0.7612 - val_loss: 0.5982 - val_acc: 0.7284
Epoch 2/10
649/649 [=====] - 0s 621us/step - loss: 0.5491 - acc: 0.7735 - val_loss: 0.5926 - val_acc: 0.7407
Epoch 3/10
649/649 [=====] - 0s 644us/step - loss: 0.5385 - acc: 0.7827 - val_loss: 0.5890 - val_acc: 0.7407
Epoch 4/10
649/649 [=====] - 0s 619us/step - loss: 0.5301 - acc: 0.7874 - val_loss: 0.5866 - val_acc: 0.7469
Epoch 5/10
649/649 [=====] - 0s 633us/step - loss: 0.5233 - acc: 0.7920 - val_loss: 0.5853 - val_acc: 0.7531
Epoch 6/10
649/649 [=====] - 0s 621us/step - loss: 0.5184 - acc: 0.7951 - val_loss: 0.5843 - val_acc: 0.7469
Epoch 7/10
649/649 [=====] - 0s 612us/step - loss: 0.5138 - acc: 0.7951 - val_loss: 0.5838 - val_acc: 0.7469
Epoch 8/10
649/649 [=====] - 0s 636us/step - loss: 0.5099 - acc: 0.7966 - val_loss: 0.5835 - val_acc: 0.7469
Epoch 9/10
649/649 [=====] - 0s 615us/step - loss: 0.5069 - acc: 0.7966 - val_loss: 0.5835 - val_acc: 0.7469
Epoch 10/10
649/649 [=====] - 0s 636us/step - loss: 0.5043 - acc: 0.7982 - val_loss: 0.5835 - val_acc: 0.7469
```




## Experiment 2: Transfer learning

### 1. BOW


```
[ ] # Our vectorized labels
y_train = np.asarray(train_labels).astype('float32')
y_test = np.asarray(test_labels).astype('float32')
```

```
[ ] from sklearn.linear_model import LogisticRegression


model = LogisticRegression()
model.fit(x_train, y_train)
```

 `sr/local/lib/python3.6/dist-packages/sklearn/linear_model/logi`  
`FutureWarning)`  
`LogisticRegression(C=1.0, class_weight=None, dual=False, fit_int`  
`intercept_scaling=1, max_iter=100, multi_class='warn',`  
`n_jobs=None, penalty='l2', random_state=None, solver='`  
`tol=0.0001, verbose=0, warm_start=False)`

```
[ ] y_pred = model.predict(x_test)
y_pred
```

 `array([0., 1., 0., ..., 0., 0., 1.], dtype=float32)`

```
[ ] from sklearn.metrics import classification_report, confusion_matr
print(classification_report(y_test, y_pred))
```



|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0          | 0.86      | 0.86   | 0.86     | 12500   |
| 1.0          | 0.86      | 0.86   | 0.86     | 12500   |
| micro avg    | 0.86      | 0.86   | 0.86     | 25000   |
| macro avg    | 0.86      | 0.86   | 0.86     | 25000   |
| weighted avg | 0.86      | 0.86   | 0.86     | 25000   |

Used Grid Search and K-fold Cross Validation

```
GridSearchCV(cv=10, error_score='raise',
             estimator=SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
                           decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
                           max_iter=-1, probability=False, random_state=None, shrinking=True,
                           tol=0.001, verbose=False),
             fit_params=None, iid=True, n_jobs=-1,
             param_grid={'C': [1, 10, 100, 1000], 'gamma': [0.1, 0.01, 1, 0.001], 'kernel': ['linear', 'rbf']},
             pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
             scoring=None, verbose=0)
```

```
[ ] y_pred = CV.predict(X_test)
     y_pred
```

```
array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1,
        1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0,
        1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1,
        1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1,
        1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1,
        1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0,
        1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0,
        1, 1, 1, 1, 1, 1, 1, 1, 1, 1])
```

```
[ ] print(metrics.accuracy_score(y_test, y_pred))
```

```
0.7914110429447853
```

```
[ ] from sklearn.metrics import classification_report, confusion_matrix
     print(classification_report(y_test, y_pred))
```

|             | precision | recall | f1-score | support |
|-------------|-----------|--------|----------|---------|
| 0           | 0.42      | 0.26   | 0.32     | 31      |
| 1           | 0.84      | 0.92   | 0.88     | 132     |
| avg / total | 0.76      | 0.79   | 0.77     | 163     |

```
[ ] confusion_matrix(y_test, y_pred)
```

```
array([[ 8, 23],
       [11, 121]], dtype=int64)
```

```
[ ] tokenize = Tokenizer()
    tokenize.fit_on_texts(transcript['text'])
    seq = tokenize.texts_to_sequences(transcript['text'])
    pad = pad_sequences(seq, maxlen = 150)
    word_idx = tokenize1.word_index
    features = pad
    features
    #features.shape
```

```
array([[ 0,  0,  0, ..., 1137, 1477, 667],
       [ 0,  0,  0, ..., 426,  51,  70],
       [ 0,  0,  0, ..., 535, 1025, 276],
       ...,
       [ 0,  0,  0, ..., 187, 1176, 709],
       [ 0,  0,  0, ..., 137, 1133, 960],
       [ 0,  0,  0, ..., 116, 1527, 474]])
```

## 2. GLOVE

Used Pretrained model

```
from keras.models import Sequential
from keras.layers import Embedding, Flatten, Dense

from keras.layers import Embedding

embedding_layer = Embedding(max_words,
                            embedding_dim,
                            embeddings_initializer=Constant(embedding_matrix),
                            input_length=maxlen,
                            trainable=False)

sequence_input = Input(shape=(maxlen,), dtype='int32')
embedded_sequences = embedding_layer(sequence_input)
x = Conv1D(128, 5, activation='relu')(embedded_sequences)
x = MaxPooling1D(5)(x)
x = Conv1D(128, 5, activation='relu', data_format='channels_first')(x)
x = MaxPooling1D(35)(x) # global max pooling
x = Flatten()(x)
x = Dense(128, activation='relu')(x)
preds = Dense(1, activation='sigmoid')(x)


model_c = Model(sequence_input, preds)
model_c.compile(loss='binary_crossentropy',
                optimizer=adam(lr=0.001),
                metrics=['acc'])


# happy learning!
history = model_c.fit(x_train, y_train, validation_data=(x_val, y_val), epochs=10, batch_size=32)
```




Train on 10000 samples, validate on 2000 samples

```
Epoch 1/10  
10000/10000 [=====] - 6s 590us/step - loss: 0.5987 - acc: 0.6656 - val_loss: 0.4501 - val_acc: 0.7950  
Epoch 2/10  
10000/10000 [=====] - 3s 292us/step - loss: 0.4087 - acc: 0.8135 - val_loss: 0.4034 - val_acc: 0.8170  
Epoch 3/10  
10000/10000 [=====] - 3s 295us/step - loss: 0.3374 - acc: 0.8544 - val_loss: 0.4723 - val_acc: 0.7865  
Epoch 4/10  
10000/10000 [=====] - 3s 294us/step - loss: 0.2791 - acc: 0.8829 - val_loss: 0.3880 - val_acc: 0.8385  
Epoch 5/10  
10000/10000 [=====] - 3s 293us/step - loss: 0.2026 - acc: 0.9185 - val_loss: 0.4081 - val_acc: 0.8385  
Epoch 6/10  
10000/10000 [=====] - 3s 299us/step - loss: 0.1208 - acc: 0.9564 - val_loss: 0.4867 - val_acc: 0.8300  
Epoch 7/10  
10000/10000 [=====] - 3s 299us/step - loss: 0.0592 - acc: 0.9819 - val_loss: 0.5700 - val_acc: 0.8330  
Epoch 8/10  
10000/10000 [=====] - 3s 303us/step - loss: 0.0291 - acc: 0.9921 - val_loss: 0.6935 - val_acc: 0.8355  
Epoch 9/10  
10000/10000 [=====] - 3s 320us/step - loss: 0.0391 - acc: 0.9864 - val_loss: 0.7809 - val_acc: 0.8190  
Epoch 10/10  
10000/10000 [=====] - 3s 331us/step - loss: 0.0507 - acc: 0.9811 - val_loss: 1.3525 - val_acc: 0.7565
```







## Word Embedding

```
from keras.models import Sequential
from keras.layers import Flatten, Dense, Dropout
from keras.optimizers import adam

model = Sequential()
model.add(Embedding(max_features, 64, input_length=maxlen))

model.add(Flatten())
model.add(Dense(32, activation='relu'))
model.add(Dropout(0.25))
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer=adam(lr=0.0001), loss='binary_crossentropy', metrics=['accuracy'])
model.summary()

train_history = model.fit(x_train, y_train,
                          epochs=10,
                          batch_size=256,
                          validation_split=0.2)
```

|                     |               |        |
|---------------------|---------------|--------|
| flatten_3 (Flatten) | (None, 12800) | 0      |
| dense_5 (Dense)     | (None, 32)    | 409632 |
| dropout_3 (Dropout) | (None, 32)    | 0      |
| dense_6 (Dense)     | (None, 1)     | 33     |

=====  
 Total params: 1,049,665  
 Trainable params: 1,049,665  
 Non-trainable params: 0


Train on 20000 samples, validate on 5000 samples


Epoch 1/10  
 20000/20000 [=====] - 4s 214us/step - loss: 0.6920 - acc: 0.5179 - val\_loss: 0.6902 - val\_acc: 0.5334  
 Epoch 2/10  
 20000/20000 [=====] - 4s 200us/step - loss: 0.6817 - acc: 0.6185 - val\_loss: 0.6823 - val\_acc: 0.5930  
 Epoch 3/10  
 20000/20000 [=====] - 4s 203us/step - loss: 0.6561 - acc: 0.6981 - val\_loss: 0.6467 - val\_acc: 0.7014  
 Epoch 4/10  
 20000/20000 [=====] - 4s 203us/step - loss: 0.5883 - acc: 0.7836 - val\_loss: 0.5638 - val\_acc: 0.7688  
 Epoch 5/10  
 20000/20000 [=====] - 4s 200us/step - loss: 0.4831 - acc: 0.8361 - val\_loss: 0.4673 - val\_acc: 0.8196  
 Epoch 6/10  
 20000/20000 [=====] - 4s 205us/step - loss: 0.3881 - acc: 0.8716 - val\_loss: 0.4013 - val\_acc: 0.8408  
 Epoch 7/10  
 20000/20000 [=====] - 4s 198us/step - loss: 0.3214 - acc: 0.8925 - val\_loss: 0.3626 - val\_acc: 0.8532  
 Epoch 8/10  
 20000/20000 [=====] - 4s 201us/step - loss: 0.2743 - acc: 0.9078 - val\_loss: 0.3371 - val\_acc: 0.8630  
 Epoch 9/10  
 20000/20000 [=====] - 4s 206us/step - loss: 0.2392 - acc: 0.9216 - val\_loss: 0.3201 - val\_acc: 0.8686  
 Epoch 10/10  
 20000/20000 [=====] - 4s 200us/step - loss: 0.2117 - acc: 0.9338 - val\_loss: 0.3086 - val\_acc: 0.8750

Test Loss: 0.30974822925567624

Test Accuracy 0.8712

```
dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
```





LSTM,

Stacked LSTM,

### STACKED LSTM

```
[ ] from keras.layers import LSTM

model_lstm = Sequential()
model_lstm.add(Embedding(max_features, 64))
model_lstm.add(LSTM(64, dropout=0.2, recurrent_dropout=0.2, return_sequences=True))
#model_lstm.add(Dropout(0.2))
model_lstm.add(LSTM(32, dropout=0.2, recurrent_dropout=0.2))

model_lstm.add(Dense(1, activation='sigmoid'))

model_lstm.compile(optimizer=adam(lr=0.0001),
                  loss='binary_crossentropy',
                  metrics=['accuracy'])

history = model_lstm.fit(x_train, y_train,
                        epochs=10,
                        batch_size=64,
                        validation_split=0.2)
```

Train on 20000 samples, validate on 5000 samples

```
Epoch 1/10
20000/20000 [=====] - 157s 8ms/step - loss: 0.6425 - acc: 0.6336 - val_loss: 0.4882 - val_acc: 0.7878
Epoch 2/10
20000/20000 [=====] - 155s 8ms/step - loss: 0.4454 - acc: 0.8076 - val_loss: 0.4223 - val_acc: 0.8152
Epoch 3/10
20000/20000 [=====] - 154s 8ms/step - loss: 0.3710 - acc: 0.8507 - val_loss: 0.3842 - val_acc: 0.8356
Epoch 4/10
20000/20000 [=====] - 157s 8ms/step - loss: 0.3313 - acc: 0.8678 - val_loss: 0.3746 - val_acc: 0.8382
Epoch 5/10
20000/20000 [=====] - 156s 8ms/step - loss: 0.3085 - acc: 0.8803 - val_loss: 0.3744 - val_acc: 0.8400
Epoch 6/10
20000/20000 [=====] - 156s 8ms/step - loss: 0.2877 - acc: 0.8913 - val_loss: 0.3756 - val_acc: 0.8412
Epoch 7/10
20000/20000 [=====] - 155s 8ms/step - loss: 0.2745 - acc: 0.8944 - val_loss: 0.3813 - val_acc: 0.8388
Epoch 8/10
20000/20000 [=====] - 156s 8ms/step - loss: 0.2582 - acc: 0.9045 - val_loss: 0.3961 - val_acc: 0.8294
Epoch 9/10
20000/20000 [=====] - 156s 8ms/step - loss: 0.2488 - acc: 0.9096 - val_loss: 0.3994 - val_acc: 0.8414
Epoch 10/10
20000/20000 [=====] - 156s 8ms/step - loss: 0.2334 - acc: 0.9148 - val_loss: 0.4042 - val_acc: 0.8398
```

Test Loss: 0.41616990515708924

Test Accuracy 0.83632

GRU(Gated Recurring Units),

RNN,



```
[ ] #Simple RNN model
    from keras.layers import Dense, SimpleRNN

    model = Sequential()
    model.add(Embedding(max_features, 32))
    model.add(SimpleRNN(32))
    model.add(Dense(1, activation='sigmoid'))

    model.compile(optimizer=adam(0.0001), loss='binary_crossentropy', metrics=['accuracy'])
    history = model.fit(x_train, y_train,
                        epochs=10,
                        batch_size=64,
                        validation_split=0.2)
```

Train on 20000 samples, validate on 5000 samples

```
Epoch 1/10
20000/20000 [=====] - 13s 632us/step - loss: 0.6946 - acc: 0.5116 - val_loss: 0.6919 - val_acc: 0.5196
Epoch 2/10
20000/20000 [=====] - 12s 607us/step - loss: 0.6795 - acc: 0.5752 - val_loss: 0.6876 - val_acc: 0.5366
Epoch 3/10
20000/20000 [=====] - 12s 607us/step - loss: 0.6026 - acc: 0.7038 - val_loss: 0.5313 - val_acc: 0.7618
Epoch 4/10
20000/20000 [=====] - 12s 611us/step - loss: 0.4479 - acc: 0.8135 - val_loss: 0.4405 - val_acc: 0.8190
Epoch 5/10
20000/20000 [=====] - 12s 618us/step - loss: 0.3451 - acc: 0.8672 - val_loss: 0.3889 - val_acc: 0.8408
Epoch 6/10
20000/20000 [=====] - 12s 623us/step - loss: 0.2835 - acc: 0.8958 - val_loss: 0.3699 - val_acc: 0.8464
Epoch 7/10
20000/20000 [=====] - 12s 621us/step - loss: 0.2360 - acc: 0.9178 - val_loss: 0.3565 - val_acc: 0.8550
Epoch 8/10
20000/20000 [=====] - 12s 617us/step - loss: 0.2023 - acc: 0.9337 - val_loss: 0.3491 - val_acc: 0.8510
Epoch 9/10
20000/20000 [=====] - 12s 611us/step - loss: 0.1745 - acc: 0.9451 - val_loss: 0.3481 - val_acc: 0.8584
Epoch 10/10
20000/20000 [=====] - 12s 611us/step - loss: 0.1448 - acc: 0.9585 - val_loss: 0.3538 - val_acc: 0.8514
```






```
import matplotlib.pyplot as plt

ax= plt.subplot()
sns.heatmap(cm, annot=True, ax = ax, fmt='g'); #annot=True to annotate cells

# labels, title and ticks
ax.set_xlabel('Predicted labels');ax.set_ylabel('True labels');
ax.set_title('Confusion Matrix');
ax.xaxis.set_ticklabels(['Negative', 'Positive']); ax.yaxis.set_ticklabels(['Positive', 'Negative']);
```



## Bi-directional RNN

We have created our own word embeddings on IMDB dataset and trained it

Used Transfer learning to predict sentiments for EDGAR datasets

## Experiment 3: Using APIs

Using the Amazon, Google, Microsoft and Watson APIs, obtain the sentiment scores for your entire dataset.

Step 1:

Fetching the initial dataset with labelled transcripts

```
[ ] with open('./final_json_data.json') as f:
    df = pd.DataFrame(json.load(f))
```

```
[ ] df.head()
```



**sentiment**

**text**

|   |          |   |
|---|----------|---|
| 0 | positive | Thank you. Good afternoon, everyone. And welco... |
| 1 | positive | As we look past Q1, we expect the channel inve... |
| 2 | neutral  | And your last question comes from the line of ... |
| 3 | negative | On the China gaming weakness, is it the slower... |
| 4 | negative | I don't know that we could tear that apart, te... |

## Step 2a : Using Amazon API - aws comprehend

```
[ ] start = time.time()
    for text in text_list:
        my_json = json.loads(json.dumps(comprehend.detect_sentiment(Text=text, LanguageCode='en'), sort_keys=True))
        amazon_sentiments_label.append(my_json['Sentiment'].lower())
        amazon_score_mixed.append(my_json['SentimentScore']['Mixed'])
        amazon_score_negative.append(my_json['SentimentScore']['Negative'])
        amazon_score_positive.append(my_json['SentimentScore']['Positive'])
        amazon_score_neutral.append(my_json['SentimentScore']['Neutral'])
    end = time.time()
    print("Time to execute", end-start)
```

Time to execute: 05.88874238053268

```
[ ] df.head()
```

|   | sentiment | text  | amazon_sentiments_label | amazon_score_mixed | amazon_score_negative | amazon_score_neutral | amazon_score_positive |
|---|-----------|---|-------------------------|--------------------|-----------------------|----------------------|-----------------------|
| 0 | positive  | Thank you. Good afternoon, everyone. And welco... | neutral                 | 0.001549           | 0.000533              | 0.750581             | 0.247337              |
| 1 | positive  | As we look past Q1, we expect the channel inve... | neutral                 | 0.009229           | 0.006072              | 0.880566             | 0.104133              |
| 2 | neutral   | And your last question comes from the line of ... | neutral                 | 0.000831           | 0.008268              | 0.980631             | 0.010270              |
| 3 | negative  | On the China gaming weakness, is it the slower... | negative                | 0.011638           | 0.953385              | 0.034877             | 0.000099              |
| 4 | negative  | I don't know that we could tear that apart, te... | positive                | 0.014957           | 0.000221              | 0.007328             | 0.977495              |

## Step 2b : Using Watson API – natural language understanding

```
[ ] import os
import json
from watson_developer_cloud import NaturalLanguageUnderstandingV1
from watson_developer_cloud.natural_language_understanding_v1 import Features, SentimentOptions
```

```
[ ] api_key = os.environ.get('IAM_ACCESS_IBM')
url = "https://gateway.watsonplatform.net/natural-language-understanding/api"
natural_language_understanding = NaturalLanguageUnderstandingV1(
    version='2018-11-16',
    iam_apikey=api_key,
    url=url
)
def ibm_sentiments(data):
# data = ("I've seen things you people wouldn't believe. Attack ships on fire off the shoulder of Orion. I watched C-t
response = natural_language_understanding.analyze(
    text=data,
    features=Features(sentiment=SentimentOptions()),
    language="en"
).get_result()
response = json.loads(json.dumps(response))
return response
```

```
[ ] start = time.time()
for text in text_list:
    my_json = ibm_sentiments(text)
    ibm_sentiments_label.append(my_json['sentiment']['document']['label'].lower())
    ibm_score.append(my_json['sentiment']['document']['score'])
end = time.time()
print("Time to execute", end-start)
```



Time to execute 626.9456360340118

| ibm_sentiments_label | ibm_score |
|----------------------|-----------|
|----------------------|-----------|

|          |          |
|----------|----------|
| positive | 0.816136 |
|----------|----------|

|          |          |
|----------|----------|
| positive | 0.558518 |
|----------|----------|

|         |          |
|---------|----------|
| neutral | 0.000000 |
|---------|----------|

|          |           |
|----------|-----------|
| negative | -0.598559 |
|----------|-----------|

|          |          |
|----------|----------|
| positive | 0.790615 |
|----------|----------|

Step 2c : Google cloud language API

```
[ ] ## This code runs on linux and not on Windows
# Imports the Google Cloud client library
from google.cloud import language
from google.cloud.language import enums
from google.cloud.language import types

# Instantiates a client
client = language.LanguageServiceClient()

def google_sentiments_api(text_in):
    # The text to analyze
    text = text_in
    document = types.Document(
        content=text,
        type=enums.Document.Type.PLAIN_TEXT)

    # Detects the sentiment of the text
    sentiment = client.analyze_sentiment(document=document).document_sentiment
    return sentiment

# print('Text: {}'.format(text))
# print('Sentiment: {}, {}'.format(sentiment.score, sentiment.magnitude))
```

The score of a document's sentiment indicates the overall emotion of a document. The magnitude of a document's sentiment indicates how much emotional content is present within the document, and this value is often proportional to the length of the document.

It is important to note that the Natural Language API indicates differences between positive and negative emotion in a document, but does not identify specific positive and negative emotions. For example, "angry" and "sad" are both considered negative emotions. However, when the Natural Language API analyzes text that is considered "angry", or text that is considered "sad", the response only indicates that the sentiment in the text is negative, not "sad" or "angry".

A document with a neutral score (around 0.0) may indicate a low-emotion document, or may indicate mixed emotions, with both high positive and negative values which cancel each out. Generally, you can use magnitude values to disambiguate these cases, as truly neutral documents will have a low magnitude value, while mixed documents will have higher magnitude values.

When comparing documents to each other (especially documents of different length), make sure to use the magnitude values to calibrate your scores, as they can help you gauge the relevant amount of emotional content.

The chart below shows some sample values and how to interpret them:

#### Sentiment Sample Values

Clearly Positive\* "score": 0.8, "magnitude": 3.0

Clearly Negative\* "score": -0.6, "magnitude": 4.0

Neutral "score": 0.1, "magnitude": 0.0

Mixed "score": 0.0, "magnitude": 4.0


\* "Clearly positive" and "clearly negative" sentiment varies for different use cases and customers.

You might find differing results for your specific scenario. We recommend that you define a threshold that works for you, and then adjust the threshold after testing and verifying the results. For example, you may define a threshold of any score over 0.25 as clearly positive, and then modify the score threshold to 0.15 after reviewing your data and results and finding that scores from 0.15-0.25 should be considered positive as well.

```
[ ] def get_label(sentiment):
    label = ""
    if sentiment.score > 0.5 and sentiment.magnitude > 1.5:
        label = "positive"
    elif sentiment.score < -0.5 and sentiment.magnitude > 1.5:
        label = "negative"
    else:
        label = "neutral"
    return label
```

```
[ ] google_sentiment_socre = []
    google_sentiment_magnitude = []
    google_sentiment_label = []
```

```
[ ] start = time.time()
    # i = 0
    for text in text_list:
        sentiment = google_sentiments_api(text)
        # print("Index" + str(i) + " Score: " + str(sentiment.score) + " Magn: " + str(sentiment.magnitude))
        google_sentiment_socre.append(sentiment.score)
        google_sentiment_magnitude.append(sentiment.magnitude)
        google_sentiment_label.append(get_label(sentiment))
    # i+=1
    end = time.time()
    print("Time to execute", end-start)
```

 Time to execute 238.47278022766113

| google_sentiment_socre | google_sentiment_magnitude | google_sentiment_label |
|------------------------|----------------------------|------------------------|
| 0.2                    | 1.1                        | positive               |
| 0.2                    | 1.4                        | positive               |
| 0.0                    | 0.0                        | neutral                |
| -0.3                   | 1.2                        | negative               |
| 0.5                    | 7.9                        | positive               |

## Step 2d: Azure text analysis API

### Preparing documents

```
document_list = []
for i, text in enumerate(text_list):
    document = {"id": str(i),
                "language": "en",
                "text": text[:5119]}
    document_list.append(document)
```

```
document_list[1]
```

```
{'id': '1',
 'language': 'en',
 'text': "As we look past Q1, we expect the channel invent
```



```
document_part_1 = {"documents" : document_list[:1000]}
```

```
document_part_2 = {"documents" : document_list[1000:]}
```

Getting sentiment score for part 1

```
[ ] start = time.time()
    headers = {"Ocp-Apim-Subscription-Key": subscription_key}
    response_1 = requests.post(sentiment_api_url, headers=headers, json=document_part_1)
    sentiments_part_1 = response_1.json()
    pprint(sentiments_part_1)
    end = time.time()
    print("Time to execute", end-start)
```

Getting sentiment score for part 2

```
[ ] start = time.time()
    headers = {"Ocp-Apim-Subscription-Key": subscription_key}
    response_2 = requests.post(sentiment_api_url, headers=headers, json=document_part_2)
    sentiments_part_2 = response_2.json()
    pprint(sentiments_part_2)
    end = time.time()
    print("Time to execute", end-start)
```

| azure_api_score | azure_api_label |
|-----------------|-----------------|
| 0.978328        | positive        |
| 0.500000        | neutral         |
| 0.500000        | neutral         |
| 0.500000        | neutral         |

Saving the file with sentiment data from all API's

```
df.to_json(path_or_buf = "final_label_json_data.json",orient='records')
```


Step 3 : Normalizing sentiment scores

- Removing Neutral sentiments

```
[ ] #Removing the neutral sentiments for our model
sentiment_df = sentiment_df[sentiment_df['sentiment'] != 'neutral']
sentiment_df.head()
```

|   | azure_api_score | google_sentiment_socre | ibm_score | amazon_sentiment_score | sentiment |
|---|-----------------|------------------------|-----------|------------------------|-----------|
| 0 | 0.978328        | 0.2                    | 0.816136  | 0.750581               | positive  |
| 1 | 0.500000        | 0.2                    | 0.558518  | 0.880566               | positive  |
| 3 | 0.500000        | -0.3                   | -0.598559 | 0.953385               | negative  |
| 4 | 0.905933        | 0.5                    | 0.790615  | 0.977495               | negative  |
| 5 | 0.904133        | 0.0                    | 0.988573  | 0.674756               | positive  |

```
[ ] sentiment_df.shape
```

 (811, 5)

- Build a model to map the output Sentiment label with the sentiment scores from all 4 apis

```
#label encode the output variable
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
sentiment_df['sentiment'] = le.fit_transform(sentiment_df['sentiment'])
```

```
sentiment_df.head()
```


|   | azure_api_score | google_sentiment_socre | ibm_score | amazon_sentiment_score | sentiment |
|---|-----------------|------------------------|-----------|------------------------|-----------|
| 0 | 0.978328        | 0.2                    | 0.816136  | 0.750581               | 1         |
| 1 | 0.500000        | 0.2                    | 0.558518  | 0.880566               | 1         |
| 3 | 0.500000        | -0.3                   | -0.598559 | 0.953385               | 0         |
| 4 | 0.905933        | 0.5                    | 0.790615  | 0.977495               | 0         |
| 5 | 0.904133        | 0.0                    | 0.988573  | 0.674756               | 1         |

- Scaling the scores

```
[ ] from sklearn.preprocessing import MinMaxScaler
```


```
[ ] X = sentiment_df.iloc[:, :-1]
    Y = sentiment_df.iloc[:, -1]
    features = X.columns.values

    features
```

 array(['azure\_api\_score', 'google\_sentiment\_socre', 'ibm\_score',  
 'amazon\_sentiment\_score'], dtype=object)

```
[ ] #Scaling all variables to a range of 0 to 1
    sc = MinMaxScaler(feature_range= (0,1))
    X = pd.DataFrame(sc.fit_transform(X), columns= features)
    #features = X.columns
```

```
[ ] features
```

 array(['azure\_api\_score', 'google\_sentiment\_socre', 'ibm\_score',  
 'amazon\_sentiment\_score'], dtype=object)

- After normalization



```
X['Avg_Norm_Sentiment_Score'] = X[['azure_api_score', 'google_sentiment_socre', 'ibm_score', 'amazon_sentiment_score']].mean(axis=1)
```

|   | azure_api_score | google_sentiment_socre | ibm_score | amazon_sentiment_score | Avg_Norm_Sentiment_Score |
|---|-----------------|------------------------|-----------|------------------------|--------------------------|
| 0 | 0.980185        | 0.5625                 | 0.903736  | 0.646323               | 0.773186                 |
| 1 | 0.468419        | 0.5625                 | 0.767878  | 0.831746               | 0.657636                 |
| 2 | 0.468419        | 0.2500                 | 0.157679  | 0.935622               | 0.452930                 |
| 3 | 0.902730        | 0.7500                 | 0.890277  | 0.970015               | 0.878256                 |
| 4 | 0.900803        | 0.4375                 | 0.994673  | 0.538159               | 0.717784                 |
| 5 | 0.823834        | 0.5000                 | 0.776448  | 0.108119               | 0.552100                 |

- Labelling w.r.t Average sentiment score

```
[ ] X.loc[ X['Avg_Norm_Sentiment_Score'] >= 0.5, 'API_Predicted_Sentiment'] = 1
X.loc[ X['Avg_Norm_Sentiment_Score'] < 0.5, 'API_Predicted_Sentiment'] = 0
X
```

|   | azure_api_score | google_sentiment_socre | ibm_score | amazon_sentiment_score | Avg_Norm_Sentiment_Score | API_Predicted_Sentiment |
|---|-----------------|------------------------|-----------|------------------------|--------------------------|-------------------------|
| 0 | 0.980185        | 0.5625                 | 0.903736  | 0.646323               | 0.773186                 | 1.0                     |
| 1 | 0.468419        | 0.5625                 | 0.767878  | 0.831746               | 0.657636                 | 1.0                     |
| 2 | 0.468419        | 0.2500                 | 0.157679  | 0.935622               | 0.452930                 | 0.0                     |
| 3 | 0.902730        | 0.7500                 | 0.890277  | 0.970015               | 0.878256                 | 1.0                     |
| 4 | 0.900803        | 0.4375                 | 0.994673  | 0.538159               | 0.717784                 | 1.0                     |

Step 4:

Getting Metrics

```
[ ] X.API_Predicted_Sentiment.value_counts()
```



```
1.0    710
0.0    101
```

Name: API\_Predicted\_Sentiment, dtype: int64

```
[ ] Y.value_counts()
```



```
1     654
0     157
```

Name: sentiment, dtype: int64

```
[ ] from sklearn.metrics import confusion_matrix
    confusion_matrix(Y, X.API_Predicted_Sentiment)
```

```
array([[ 41, 116],
       [ 60, 594]], dtype=int64)
```

```
[ ] from sklearn.metrics import accuracy_score
    print(accuracy_score(Y, X.API_Predicted_Sentiment))
```

```
0.782983970406905
```

```
[ ] from sklearn.metrics import classification_report
    print(classification_report(Y, X.API_Predicted_Sentiment))
```

|             | precision | recall | f1-score | support |
|-------------|-----------|--------|----------|---------|
| 0           | 0.41      | 0.26   | 0.32     | 157     |
| 1           | 0.84      | 0.91   | 0.87     | 654     |
| avg / total | 0.75      | 0.78   | 0.76     | 811     |


Plotting the Confusion matrix in Heatmap

```
[ ] import matplotlib.pyplot as plt
    import seaborn as sns

    ax= plt.subplot()
    #cm = classification_report(y_test, y_pred_lstm)
    sns.heatmap(confusion_matrix(Y, X.API_Predicted_Sentiment), annot=True, ax = ax, fmt='g'); #annot=True to annotate cells

    # labels, title and ticks
    ax.set_xlabel('Predicted labels');ax.set_ylabel('True labels');
    ax.set_title('Confusion Matrix');
    ax.xaxis.set_ticklabels(['Negative', 'Positive'])
    ax.yaxis.set_ticklabels(['Positive', 'Negative'])
```

```
[Text(0,0.5,'Negative'), Text(0,1.5,'Positive')]
```



#### Experiment 4: Ensemble learning using AutoML

In order to map the sentiment scores from all 4 API's to the sentiment's we labelled manually, we created a FC Neural network

```
[ ] from keras.layers import Dense, Dropout
    from keras.models import Sequential

    model = Sequential()

[ ] model.add(Dense(32, activation='relu', kernel_initializer='uniform', input_shape=(4,)))
    model.add(Dropout(0.15))
    model.add(Dense(16, activation='relu', kernel_initializer='uniform'))
    model.add(Dropout(0.2))
    model.add(Dense(1, activation='sigmoid', kernel_initializer='uniform'))

    model.summary()

    model.compile(optimizer='adam', loss= 'binary_crossentropy', metrics=['accuracy'])
```





| Layer (type)            | Output Shape | Param # |
|-------------------------|--------------|---------|
| dense_25 (Dense)        | (None, 32)   | 160     |
| dropout_17 (Dropout)    | (None, 32)   | 0       |
| dense_26 (Dense)        | (None, 16)   | 528     |
| dropout_18 (Dropout)    | (None, 16)   | 0       |
| dense_27 (Dense)        | (None, 1)    | 17      |
| =====                   |              |         |
| Total params: 705       |              |         |
| Trainable params: 705   |              |         |
| Non-trainable params: 0 |              |         |

Followed by this we tried TPOT, AutoSKlearn , H2O.ai in order to utilize autoML for mapping the same

Test Loss: 0.41255556056104553

Test Accuracy 0.8098159509202454

```
dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
```



We did hyperparameter tuning for the model so that the loss converges better

|                      |            |      |
|----------------------|------------|------|
| dense_28 (Dense)     | (None, 32) | 160  |
| dropout_19 (Dropout) | (None, 32) | 0    |
| dense_29 (Dense)     | (None, 32) | 1056 |
| dropout_20 (Dropout) | (None, 32) | 0    |
| dense_30 (Dense)     | (None, 1)  | 33   |
| dense_31 (Dense)     | (None, 32) | 64   |
| dropout_21 (Dropout) | (None, 32) | 0    |
| dense_32 (Dense)     | (None, 32) | 1056 |
| dropout_22 (Dropout) | (None, 32) | 0    |
| dense_33 (Dense)     | (None, 1)  | 33   |
| dense_34 (Dense)     | (None, 32) | 64   |
| dropout_23 (Dropout) | (None, 32) | 0    |
| dense_35 (Dense)     | (None, 32) | 1056 |
| dropout_24 (Dropout) | (None, 32) | 0    |
| dense_36 (Dense)     | (None, 1)  | 33   |

```



=====
C:\Users\nikhi\Anaconda3\lib\site-packages\keras\engine\training.py:
'Discrepancy between trainable weights and collected trainable'
Total params: 2,402
Trainable params: 2,402
Non-trainable params: 0

```

Test Loss: 0.6833624390005334

Test Accuracy 0.8098159509202454

```
dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
```



Since the model is trained on very few negative data, the model doesn't predict the negative sentiments at all and we find 0 True Negatives as a result


```
from sklearn.metrics import confusion_matrix, classification_report
confusion_matrix(Y_test, y_pred_class)
```

```
array([[ 0, 31],
       [ 0, 132]], dtype=int64)
```

```
ax= plt.subplot()
#cm = classification_report(y_test, y_pred_lstm)
sns.heatmap(confusion_matrix(Y_test, y_pred_class), annot=True, ax = ax, fmt='g');

# labels, title and ticks
ax.set_xlabel('Predicted labels');ax.set_ylabel('True labels');
ax.set_title('Confusion Matrix');
ax.xaxis.set_ticklabels(['Negative', 'Positive'])
ax.yaxis.set_ticklabels(['Negative', 'Positive'])
```

```
[Text(0,0.5,'Negative'), Text(0,1.5,'Positive')]
```



## AUTO ML

Implementing TPOT to get a model with tuned hyperparameters



```
[ ] from tpot import TPOTClassifier
    #from tpot import TPOTRegressor

    tpot = TPOTClassifier(generations=5,verbosity=2)
```

```
[ ] tpot.fit(X_train,Y_train)
```

```

HBox(children=(IntProgress(value=0, description='Optimization Progress', max=600), HTML(value='')))
Generation 1 - Current best internal CV score: 0.8009970667380429
Generation 2 - Current best internal CV score: 0.8009970667380429
Generation 3 - Current best internal CV score: 0.8009970667380429
Generation 4 - Current best internal CV score: 0.8009970667380429
Generation 5 - Current best internal CV score: 0.8009970667380429

Best pipeline: ExtraTreesClassifier(input_matrix, bootstrap=True, criterion=entropy, max_features=0.8500000000000001,
TPOTClassifier(config_dict=None, crossover_rate=0.1, cv=5,
               disable_update_check=False, early_stop=None, generations=5,
               max_eval_time_mins=5, max_time_mins=None, memory=None,
               mutation_rate=0.9, n_jobs=1, offspring_size=None,
               periodic_checkpoint_folder=None, population_size=100,
               random_state=None, scoring=None, subsample=1.0, use_dask=False,
               verbosity=2, warm_start=False))

```

## Output of TPOT:

```
[ ] tpot.fitted_pipeline_
```

```

Pipeline(memory=None,
         steps=[('extratreesclassifier', ExtraTreesClassifier(bootstrap=True, class_weight=None, criterion='entropy',
max_depth=None, max_features=0.8500000000000001,
max_leaf_nodes=None, min_impurity_decrease=0.0,
min_impurity_split=None, min_samples_leaf=19,
min_samples_split=11, min_weight_fraction_leaf=0.0,
n_estimators=100, n_jobs=1, oob_score=False, random_state=None,
verbose=0, warm_start=False))])

```

So the best pipeline as per TPOT is the below -

Best pipeline: KNeighborsClassifier(input\_matrix, n\_neighbors=90, p=1, weights=distance)

```
[ ] # !conda install -c anaconda py-xgboost
```

```
[ ] tpot.score(X_test,
               Y_test)
```

```
0.8226600985221675
```

## Using H2O for creating a complete pipeline

```
[ ] import h2o
    from h2o.automl import H2OAutoML
    h2o.init()
```

Checking whether there is an H2O instance running at <http://localhost:54321>..... not found.  
 Attempting to start a local H2O server...  
 ; Java HotSpot(TM) 64-Bit Server VM (build 25.191-b12, mixed mode)  
 Starting server from C:\Users\nikhi\Anaconda3\lib\site-packages\h2o\backend\bin\h2o.jar  
 Ice root: C:\Users\nikhi\AppData\Local\Temp\tmp3fv192bt  
 JVM stdout: C:\Users\nikhi\AppData\Local\Temp\tmp3fv192bt\h2o\_nikhi\_started\_from\_python.out  
 JVM stderr: C:\Users\nikhi\AppData\Local\Temp\tmp3fv192bt\h2o\_nikhi\_started\_from\_python.err  
 Server is running at <http://127.0.0.1:54321>  
 Connecting to H2O server at <http://127.0.0.1:54321>... successful.  
 H2O cluster uptime: 01 secs  
 H2O cluster timezone: America/New\_York  
 H2O data parsing timezone: UTC  
 H2O cluster version: 3.22.1.2  
 H2O cluster version age: 2 months and 1 day  
 H2O cluster name: H2O\_from\_python\_nikhi\_dvh1lf  
 H2O cluster total nodes: 1  
 H2O cluster free memory: 3.523 Gb  
 H2O cluster total cores: 8  
 H2O cluster allowed cores: 8  
 H2O cluster status: accepting new members, healthy  
 H2O connection url: http://127.0.0.1:54321  
 H2O connection proxy: None  
 H2O internal security: False  
 H2O API Extensions: Algos, AutoML, Core V3, Core V4  
 Python version: 3.6.5 final

## Description about data

```
[ ] df.describe()
```

Rows:811  
 Cols:6

|         | C1                 | azure_api_score     | google_sentiment_socre | ibm_score          | amazon_sentiment_score | sentiment           |
|---------|--------------------|---------------------|------------------------|--------------------|------------------------|---------------------|
| type    | int                | real                | real                   | real               | real                   | int                 |
| mins    | 0.0                | 0.0621877611        | -0.6999999881          | -0.897554          | 0.2974953949           | 0.0                 |
| mean    | 828.5326757090012  | 0.6427028278373615  | 0.21418002826979035    | 0.5957383181257702 | 0.7602460733112212     | 0.8064118372379778  |
| maxs    | 1642.0             | 0.9968479276        | 0.8999999762           | 0.998674           | 0.9985154271           | 1.0                 |
| sigma   | 455.07252153575007 | 0.22324971364639776 | 0.24069582961015862    | 0.4168205697793954 | 0.17980453185630274    | 0.39535366015815204 |
| zeros   | 1                  | 0                   | 187                    | 18                 | 0                      | 157                 |
| missing | 0                  | 0                   | 0                      | 0                  | 0                      | 0                   |
| 0       | 0.0                | 0.9783278108        | 0.200000003            | 0.816136           | 0.7505810857           | 1.0                 |
| 1       | 1.0                | 0.5                 | 0.200000003            | 0.558518           | 0.8805660009           | 1.0                 |
| 2       | 3.0                | 0.5                 | -0.3000000119          | -0.598559          | 0.9533853531           | 0.0                 |
| 3       | 4.0                | 0.9059334397        | 0.5                    | 0.790615           | 0.9774951339           | 0.0                 |
| 4       | 5.0                | 0.9041327238        | 0.0                    | 0.988573           | 0.6747556925           | 1.0                 |
| 5       | 6.0                | 0.8321921229        | 0.1000000015           | 0.574769           | 0.3732891977           | 0.0                 |
| 6       | 8.0                | 0.5                 | 0.400000006            | 0.884016           | 0.7173274755           | 1.0                 |
| 7       | 9.0                | 0.7486802936        | 0.6000000238           | 0.59894            | 0.6371069551           | 1.0                 |
| 8       | 10.0               | 0.9522012472        | 0.400000006            | 0.875206           | 0.9451110363           | 1.0                 |
| 9       | 12.0               | 0.5                 | 0.200000003            | 0.569009           | 0.6385424137           | 0.0                 |

Training 10 models and getting there metrics score as follows:

| predict | p0        | p1       |
|---------|-----------|----------|
| 1       | 0.0591668 | 0.940833 |
| 1       | 0.0715695 | 0.928431 |
| 1       | 0.102093  | 0.897907 |
| 1       | 0.0987179 | 0.901282 |
| 1       | 0.0680783 | 0.931922 |
| 1       | 0.0665929 | 0.933407 |
| 1       | 0.0667654 | 0.933235 |
| 1       | 0.0626903 | 0.93731  |
| 1       | 0.0310117 | 0.968988 |
| 1       | 0.101915  | 0.898085 |

## AutoSKlearn

### Facing issue with autosklearn package

The issue is still open and the link to it is down below

```
[ ] import autosklearn.classification
import sklearn.model_selection
import sklearn.datasets
import sklearn.metrics
from sklearn.model_selection import cross_val_score, train_test_split

[ ] X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=42)
automl = autosklearn.classification.AutoSklearnClassifier()
automl.fit(X_train, Y_train)
y_hat = automl.predict(X_test)
print("Accuracy score", sklearn.metrics.accuracy_score(Y_test, y_hat))
```

me/jai/miniconda3/lib/python3.7/site-packages/autosklearn/evaluation/train\_evaluator.py:197: RuntimeWarning: Mean of empty slice  
/train\_pred = np.nanmean(Y\_train\_pred\_full, axis=0)  
ARNING] [2019-03-21 08:38:02,092:EnsembleBuilder(1):9af0d6fbec1b15d1b06c1bc99a3f82d0] No models better than random - using Dummy Score!  
ARNING] [2019-03-21 08:38:02,100:EnsembleBuilder(1):9af0d6fbec1b15d1b06c1bc99a3f82d0] No models better than random - using Dummy Score!  
me/jai/miniconda3/lib/python3.7/site-packages/autosklearn/evaluation/train\_evaluator.py:197: RuntimeWarning: Mean of empty slice  
/train\_pred = np.nanmean(Y\_train\_pred\_full, axis=0)  
me/jai/miniconda3/lib/python3.7/site-packages/autosklearn/evaluation/train\_evaluator.py:197: RuntimeWarning: Mean of empty slice  
/train\_pred = np.nanmean(Y\_train\_pred\_full, axis=0)  
me/jai/miniconda3/lib/python3.7/site-packages/autosklearn/evaluation/train\_evaluator.py:197: RuntimeWarning: Mean of empty slice  
/train\_pred = np.nanmean(Y\_train\_pred\_full, axis=0)

<https://github.com/automl/auto-sklearn/issues/520>